

In [1]:

```
#Import required libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

Acquire Housing dataset

In [2]:

```
#Read housing dataset
df_housing_dataset = pd.read_csv('F:\SML\housing.csv')
```

In [3]:

```
df_housing_dataset.head()
```

Out[3]:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	populati
0	-122.23	37.88	41	880	129.0	322
1	-122.22	37.86	21	7099	1106.0	2401
2	-122.24	37.85	52	1467	190.0	496
3	-122.25	37.85	52	1274	235.0	558
4	-122.25	37.85	52	1627	280.0	565

In [4]:

```
df_housing_dataset.columns
```

Out[4]:

```
Index(['longitude', 'latitude', 'housing_median_age', 'total_rooms',
      'total_bedrooms', 'population', 'households', 'median_income',
      'ocean_proximity', 'median_house_value'],
      dtype='object')
```

In [5]:

```
#Check shape of entire dataset
df_housing_dataset.shape
```

Out[5]:

```
(20640, 10)
```

In [6]:

```
df_housing_dataset.describe()
```

Out[6]:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms
count	20640.000000	20640.000000	20640.000000	20640.000000	20433.000000
mean	-119.569704	35.631861	28.639486	2635.763081	537.870553
std	2.003532	2.135952	12.585558	2181.615252	421.385070
min	-124.350000	32.540000	1.000000	2.000000	1.000000
25%	-121.800000	33.930000	18.000000	1447.750000	296.000000
50%	-118.490000	34.260000	29.000000	2127.000000	435.000000
75%	-118.010000	37.710000	37.000000	3148.000000	647.000000
max	-114.310000	41.950000	52.000000	39320.000000	6445.000000

Visualize data to understand the relationship among variables

In [7]:

```
corr = df_housing_dataset.corr()
```

In [8]:


```
df_housing_dataset.corr()
```

Out[8]:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms
longitude	1.000000	-0.924664	-0.108197	0.044568	0.069608
latitude	-0.924664	1.000000	0.011173	-0.036100	-0.066983
housing_median_age	-0.108197	0.011173	1.000000	-0.361262	-0.320451
total_rooms	0.044568	-0.036100	-0.361262	1.000000	0.930380
total_bedrooms	0.069608	-0.066983	-0.320451	0.930380	1.000000
population	0.099773	-0.108785	-0.296244	0.857126	0.884627
households	0.055310	-0.071035	-0.302916	0.918484	0.918484
median_income	-0.015176	-0.079809	-0.119034	0.198050	-0.015176
median_house_value	-0.045967	-0.144160	0.105623	0.134153	0.045967


In [9]:

```
#Seaborn heatmap to view correlations between features in dataset  
#Median income has a positive correlation against median house value  
plt.figure(figsize=(8,6))  
plt.heatmap = sns.heatmap(corr)
```



In [10]:

```
#Scatter plot of median income with median house value
plt.figure(figsize=(10,7))
plt.scatter(df_housing_dataset['median_income'],df_housing_dataset['median_house_value'])
plt.title='Scatter plot to correlate median income vs median house value'
plt.xlabel('Median Income')
plt.ylabel('Medain House Value')
plt.show()
```



In [11]:

```
df_housing_dataset.ocean_proximity.unique()
```

Out[11]:

```
array(['NEAR BAY', '<1H OCEAN', 'INLAND', 'NEAR OCEAN', 'ISLAND'],
      dtype=object)
```

In [12]:

```
df_housing_dataset.ocean_proximity.isnull().sum()
```

Out[12]:

0

In [13]:

```
#Slice dataset and store independent and dependent variables
```

```
X = df_housing_dataset.iloc[:, :-1].values
```

```
y = df_housing_dataset.iloc[:, 9].values
```

```
print (X,y)
```

```
[[-122.23  37.88  41 ... 126  8.3252 'NEAR BAY']
```

```
 [-122.22  37.86  21 ... 1138  8.3014 'NEAR BAY']
```

```
 [-122.24  37.85  52 ... 177  7.2574 'NEAR BAY']
```

```
...
```

```
 [-121.22  39.43  17 ... 433  1.7 'INLAND']
```

```
 [-121.32  39.43  18 ... 349  1.8672 'INLAND']
```

```
 [-121.24  39.37  16 ... 530  2.3886 'INLAND']] [452600 358500 352100 ... 92  
300 84700 89400]
```

In [14]:

```
#Label Encode ocean proximity column
```

```
from sklearn.preprocessing import LabelEncoder
```

```
ocean_proximity_labelencoder = LabelEncoder()
```

```
X[:,8] = ocean_proximity_labelencoder.fit_transform(X[:,8])
```

In [15]:

```
X[:,8]
```

Out[15]:

```
array([3, 3, 3, ..., 1, 1, 1], dtype=object)
```

In [16]:

```
#Correlation between Ocean proximity and Median house value
```

```
corr1 = np.corrcoef(X[:,8].astype('float64'),y.astype('float64'))
```

```
print(corr1)
```

```
[[1.          0.08175023]
```


```
 [0.08175023 1.          ]]
```

In [17]:

```
#Seaborn pairplot of median income vs median house value with hue as Ocean Proximity
#hue="ocean_proximity"
sns.pairplot(df_housing_dataset, size=4 ,
             vars=["median_income", "median_house_value"], hue="ocean_proximity")
```

Out[17]:

<seaborn.axisgrid.PairGrid at 0x22c5ac79dd8>



Handle missing values

In [18]:

```
df_housing_dataset.isnull().sum()
```

Out[18]:

```
longitude          0
latitude           0
housing_median_age  0
total_rooms         0
total_bedrooms     207
population          0
households          0
median_income       0
ocean_proximity    0
median_house_value  0
dtype: int64
```

In [19]:

```
df_X = pd.DataFrame(X)
```

In [20]:

```
df_X.isnull().sum()
```

Out[20]:

```
0      0
1      0
2      0
3      0
4     207
5      0
6      0
7      0
8      0
dtype: int64
```

In [24]:

```
#####
# Handle the missing values, we can see that in dataset there are some missing
# values, we will use strategy to impute mean of column values in these places
#####>>>
from sklearn.impute import SimpleImputer

#from sklearn.preprocessing import Imputer
# First create an Imputer
missingValueImputer = Imputer (missing_values = 'NaN', strategy = 'mean',
                                axis = 0)
# Set which columns imputer should perform
missingValueImputer = missingValueImputer.fit (X[:,4:5])
# update values of X with new values
X[:,4:5] = missingValueImputer.transform(X[:,4:5])
```

```
C:\Users\SaiRam\Anaconda3\lib\site-packages\sklearn\utils\deprecation.py:5
8: DeprecationWarning: Class Imputer is deprecated; Imputer was deprecated
in version 0.20 and will be removed in 0.22. Import impute.SimpleImputer f
rom sklearn instead.
   warnings.warn(msg, category=DeprecationWarning)
```

In [25]:

```
#Notice missing values in total bedrooms column have been imputed with mean of total be
drooms
df_X = pd.DataFrame(X)
df_X.isnull().sum()
```

Out[25]:

```
0    0
1    0
2    0
3    0
4    0
5    0
6    0
7    0
8    0
dtype: int64
```

Principal Component Analysis

In [26]:

```
X.shape
```

Out[26]:

```
(20640, 9)
```


In [27]:

```
#Feature Scaling
#from sklearn.preprocessing import StandardScaler
#stdsclr = StandardScaler()
#X_std = stdsclr.fit_transform(X)
from sklearn.preprocessing import StandardScaler
stdsclr = StandardScaler()
(stdsclr.fit(X))
X_std = stdsclr.fit_transform(X)
```

```
C:\Users\SaiRam\Anaconda3\lib\site-packages\sklearn\utils\validation.py:59
5: DataConversionWarning: Data with input dtype object was converted to fl
oat64 by StandardScaler.
warnings.warn(msg, DataConversionWarning)
C:\Users\SaiRam\Anaconda3\lib\site-packages\sklearn\utils\validation.py:59
5: DataConversionWarning: Data with input dtype object was converted to fl
oat64 by StandardScaler.
warnings.warn(msg, DataConversionWarning)
C:\Users\SaiRam\Anaconda3\lib\site-packages\sklearn\utils\validation.py:59
5: DataConversionWarning: Data with input dtype object was converted to fl
oat64 by StandardScaler.
warnings.warn(msg, DataConversionWarning)
```

In [28]:

```
#PCA
from sklearn.decomposition.pca import PCA
PCA = PCA(n_components=6)
principal_components = PCA.fit_transform(X_std)
```

In [29]:

```
principal_components
```

Out[29]:

```
array([[ -2.15719994,  1.70225453,  1.8547863 ,  1.70382331,  0.70467378,
         0.14044674],
       [ 2.87263151,  2.3047156 ,  1.9782144 ,  1.45640433,  0.22588098,
        -0.38227627],
       [-2.03519184,  1.79186986,  0.9709968 ,  1.89443788,  1.26251199,
         0.04699636],
       ...,
       [-0.45515873,  1.6384034 , -0.31193923, -1.51947927, -0.81951544,
        -0.24827694],
       [-0.86707319,  1.61941222, -0.20800024, -1.4886554 , -0.79020992,
        -0.28420985],
       [ 0.13316801,  1.6848017 , -0.05224996, -1.36842899, -0.68749535,
        -0.17045861]])
```

In [30]:

```
#Cal the cumulative proportion of var explained by each component
PCA.explained_variance_ratio_
```

Out[30]:

```
array([0.43398133, 0.22466942, 0.11957182, 0.10398544, 0.08531692,
       0.01648695])
```

In [31]:

```
df_X = pd.DataFrame(X)
print(df_X.columns)
```

RangeIndex(start=0, stop=9, step=1)

In [32]:

```
df_housing_dataset.head()
```

Out[32]:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	populati
0	-122.23	37.88	41	880	129.0	322
1	-122.22	37.86	21	7099	1106.0	2401
2	-122.24	37.85	52	1467	190.0	496
3	-122.25	37.85	52	1274	235.0	558
4	-122.25	37.85	52	1627	280.0	565

In [33]:

```
# Dump components relations with features: This gives us the picture of how features are related to components
print(pd.DataFrame(PCA.components_, columns=df_X.columns, index = ['PC-1', 'PC-2', 'PC-3', 'PC-4', 'PC-5', 'PC-6']))
```

	0	1	2	3	4	5	6
PC-1	0.081446	-0.077765	-0.219732	0.482987	0.488518	0.471762	0.490642
PC-2	-0.670071	0.655264	0.033190	0.084062	0.072089	0.031852	0.074866
PC-3	-0.089342	0.065996	-0.428611	0.085889	-0.120442	-0.114825	-0.113064
PC-4	0.110276	-0.277884	0.419471	0.082480	0.029807	0.002983	0.041821
PC-5	-0.140912	0.061118	0.762079	0.085413	0.046079	0.096782	0.078822
PC-6	-0.113470	-0.073868	-0.042409	-0.313566	-0.391694	0.841691	-0.123976

	7	8
PC-1	0.045539	-0.041798
PC-2	-0.032873	0.317125
PC-3	0.856744	-0.148639
PC-4	0.377072	0.763565
PC-5	0.290296	-0.535139
PC-6	0.052332	0.039623

In [34]:

```
principal_components.shape
```

Out[34]:

```
(20640, 6)
```

Machine Learning Model Selection and Training

In [35]:

```
#Let's check our target label  
y
```

Out[35]:

```
array([452600, 358500, 352100, ..., 92300, 84700, 89400], dtype=int64)
```

In [36]:

```
#Split Dataset for model training and testing [80/20 split]  
#from sklearn.cross_validation import train_test_split  
from sklearn.model_selection import train_test_split  
X_train,X_test,y_train,y_test = train_test_split(principal_components,y, test_size=0.1,  
random_state=1)
```

In [37]:

```
X_train.shape
```

Out[37]:

```
(18576, 6)
```

Linear Regression ML Model

In [38]:

```
#Linear Regression Model  
from sklearn.linear_model import LinearRegression  
linReg = LinearRegression()  
linReg.fit(X_train,y_train)
```

Out[38]:

```
LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,  
normalize=False)
```

In [39]:

```
linReg.predict(X_test)
```

Out[39]:

```
array([216366.32806176, 141489.28480139, 236970.65976977, ...,  
407406.64357423, 208935.87886274, 199176.10996327])
```

In [40]:

```
#Quick check accuracy of the model
score = linReg.score(X_train,y_train)
print(score)
```

0.5415486643334289

In [41]:

```
#Quick check accuracy of the model
score = linReg.score(X_test,y_test)
print(score)
```

0.5350362554208468

In [42]:

```
linreg_predictions = linReg.predict(X_test)
```

In [43]:

```
from sklearn.metrics import mean_squared_error
np.sqrt(mean_squared_error(y_test,linreg_predictions))
#RMSE below
```

Out[43]:

79072.94760163683

Decision Tree ML model

In [44]:

```
#Train with DT model
from sklearn.tree import DecisionTreeRegressor
DTRegressor = DecisionTreeRegressor(max_depth=9, min_samples_split=5)
DTRegressor.fit(X_train,y_train)
```

Out[44]:

```
DecisionTreeRegressor(criterion='mse', max_depth=9, max_features=None,
                      max_leaf_nodes=None, min_impurity_decrease=0.0,
                      min_impurity_split=None, min_samples_leaf=1,
                      min_samples_split=5, min_weight_fraction_leaf=0.0,
                      presort=False, random_state=None, splitter='best')
```

In [45]:

```
#Quick check accuracy of the model
score = DTRegressor.score(X_train,y_train)
print(score)
```

0.7538343085921241

In [46]:

```
#Quick check accuracy of the model  
score = DTRegressor.score(X_test,y_test)  
print(score)
```

0.6347818324506117

In [47]:

```
DTR_predictions = DTRegressor.predict(X_test)
```

In [48]:

```
from sklearn.metrics import mean_squared_error  
np.sqrt(mean_squared_error(y_test,DTR_predictions))  
#RMSE below
```

Out[48]:

70080.07690286402

Random Forest ML Model - Model prediction accuracy is good compared to LR and DT models

In [49]:

```
#Declare hyper parameters to tune RF model  
hyperparameters = { 'randomforestregressor__max_features' : ['auto', 'sqrt', 'log2'],  
                    'randomforestregressor__max_depth': [None, 5, 3, 1],  
                    'randomforestregressor__min_samples_split': [2, 5],  
                    'randomforestregressor__min_samples_leaf': [10, 5]}
```

In [50]:

```
#Make a Random forest pipeline  
from sklearn.ensemble import RandomForestRegressor  
from sklearn.pipeline import make_pipeline  
pipeline = make_pipeline(RandomForestRegressor(n_estimators=50))
```

In [52]:

```
#Cross Validation to find best parameters
#from sklearn.grid_search import GridSearchCV
from sklearn.model_selection import GridSearchCV
clf = GridSearchCV(pipeline, hyperparameters, cv=10)
# Fit and tune model
clf.fit(X_train, y_train)
```

Out[52]:

```
GridSearchCV(cv=10, error_score='raise-deprecating',
             estimator=Pipeline(memory=None,
                                 steps=[('randomforestregressor', RandomForestRegressor(bootstrap=True,
criterion='mse', max_depth=None,
max_features='auto', max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, n_estimators=50, n_jobs=None,
oob_score=False, random_state=None, verbose=0, warm_start=False
e))),
             fit_params=None, iid='warn', n_jobs=None,
             param_grid={'randomforestregressor__max_features': ['auto', 'sqrt',
'log2'], 'randomforestregressor__max_depth': [None, 5, 3, 1], 'randomfores
tregressor__min_samples_split': [2, 5], 'randomforestregressor__min_sample
s_leaf': [10, 5]},
             pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
             scoring=None, verbose=0)
```

In [53]:

```
clf.best_params_
```

Out[53]:

```
{'randomforestregressor__max_depth': None,
 'randomforestregressor__max_features': 'auto',
 'randomforestregressor__min_samples_leaf': 5,
 'randomforestregressor__min_samples_split': 5}
```

In [54]:

```
clf.best_score_
```

Out[54]:

```
0.7427108313947447
```

In [55]:

```
clf.best_estimator_
```

Out[55]:

```
Pipeline(memory=None,
          steps=[('randomforestregressor', RandomForestRegressor(bootstrap=True,
          criterion='mse', max_depth=None,
          max_features='auto', max_leaf_nodes=None,
          min_impurity_decrease=0.0, min_impurity_split=None,
          min_samples_leaf=5, min_samples_split=5,
          min_weight_fraction_leaf=0.0, n_estimators=50, n_jobs=None,
          oob_score=False, random_state=None, verbose=0, warm_start=False
          e))])
```

In [56]:

```
#Quick check accuracy of the model after CV
score = clf.score(X_train,y_train)
print(score)
```

```
0.8789239810104083
```

In [57]:

```
#Quick check accuracy of the model
score = clf.score(X_test,y_test)
print(score)
```

```
0.7389279737958411
```

In [58]:

```
X_test
```

Out[58]:

```
array([[ 0.165356 , -1.06409846, -0.64909599, -0.10464544,  0.86158045,
        -0.52274848],
       [-0.62981174,  0.38872249, -1.02298362, -0.37742015,  0.26363882,
        -0.16050481],
       [-0.94940543,  1.98117187, -0.50598625,  1.11304819,  0.5866817 ,
        0.43237097],
       ...,
       [ 9.94353045, -0.33380172, -0.03629273,  0.54202254, -0.30348569,
        -4.77201298],
       [-1.87167268, -1.53793658, -0.0138332 , -0.23525572,  0.41688408,
        -0.09706565],
       [-1.53844573,  1.2624476 ,  0.47595115, -0.61976029,  0.10913802,
        -0.04103344]])
```

In [59]:

```
RF_predictions = clf.predict(X_test)
```

In [60]:

```
from sklearn.metrics import mean_squared_error
np.sqrt(mean_squared_error(y_test,RF_predictions))
#RMSE below for the 10% test set [unseen data]
```

Out[60]:

59251.391364186624

ML Model Training with only Median income feature to predict housing value

In [61]:

```
#Let's train the model only with median income and check how model behaves
X
```

Out[61]:

```
array([[ -122.23, 37.88, 41, ..., 126, 8.3252, 3],
       [ -122.22, 37.86, 21, ..., 1138, 8.3014, 3],
       [ -122.24, 37.85, 52, ..., 177, 7.2574, 3],
       ...,
       [ -121.22, 39.43, 17, ..., 433, 1.7, 1],
       [ -121.32, 39.43, 18, ..., 349, 1.8672, 1],
       [ -121.24, 39.37, 16, ..., 530, 2.3886, 1]], dtype=object)
```

In [62]:

```
df_X_final = pd.DataFrame(X)
df_X_final.head()
```

Out[62]:

	0	1	2	3	4	5	6	7	8
0	-122.23	37.88	41	880	129	322	126	8.3252	3
1	-122.22	37.86	21	7099	1106	2401	1138	8.3014	3
2	-122.24	37.85	52	1467	190	496	177	7.2574	3
3	-122.25	37.85	52	1274	235	558	219	5.6431	3
4	-122.25	37.85	52	1627	280	565	259	3.8462	3

In [63]:

```
X = np.delete(X,[0,1,2,3,4,5,6,8],axis=1)
```


In [64]:

```
X
```

Out[64]:

```
array([[8.3252],  
       [8.3014],  
       [7.2574],  
       ...,  
       [1.7],  
       [1.8672],  
       [2.3886]], dtype=object)
```

In [77]:

```
#Split Dataset for model training and testing [80/20 split]  
from sklearn.model_selection import train_test_split  
X_train,X_test,y_train,y_test = train_test_split(X,y, test_size=1/4,random_state=0)
```

In [78]:

```
X_train.shape
```

Out[78]:

```
(15480, 1)
```

In [79]:

```
#Feature Scaling  
from sklearn.preprocessing import StandardScaler  
stdsclr = StandardScaler()  
X_train_std = stdsclr.fit_transform(X_train)
```

```
C:\Users\SaiRam\Anaconda3\lib\site-packages\sklearn\utils\validation.py:59  
5: DataConversionWarning: Data with input dtype object was converted to fl  
oat64 by StandardScaler.  
  warnings.warn(msg, DataConversionWarning)  
C:\Users\SaiRam\Anaconda3\lib\site-packages\sklearn\utils\validation.py:59  
5: DataConversionWarning: Data with input dtype object was converted to fl  
oat64 by StandardScaler.  
  warnings.warn(msg, DataConversionWarning)
```

In [80]:

```
#Feature Scaling  
X_test_std = stdsclr.fit_transform(X_test)
```

```
C:\Users\SaiRam\Anaconda3\lib\site-packages\sklearn\utils\validation.py:59  
5: DataConversionWarning: Data with input dtype object was converted to fl  
oat64 by StandardScaler.  
  warnings.warn(msg, DataConversionWarning)  
C:\Users\SaiRam\Anaconda3\lib\site-packages\sklearn\utils\validation.py:59  
5: DataConversionWarning: Data with input dtype object was converted to fl  
oat64 by StandardScaler.  
  warnings.warn(msg, DataConversionWarning)
```

In [81]:

```
#Linear Regression Model
from sklearn.linear_model import LinearRegression
linReg1 = LinearRegression()
linReg1.fit(X_train_std,y_train)
```

Out[81]:

```
LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,
normalize=False)
```

In [82]:

```
#Quick check accuracy of the model
score = linReg1.score(X_train_std,y_train)
print(score)
```

0.48061930819884535


In [83]:

```
#Quick check accuracy of the model
score = linReg1.score(X_test_std,y_test)
print(score)
```

0.45147717106069024

In [84]:

```
#=====
# Visualize the linear regressor algo outcome
#=====
# Visualising the Regression results
plt.scatter(X_test, y_test, color = 'red')
plt.plot(X_test, linReg1.predict(X_test), color = 'blue')
#plt.title('Median House Price Prediction')
plt.xlabel('Median Income')
plt.ylabel('Median House Price')
plt.show()
```



In [85]:

```
#=====
# Fitting the Polynomial Regression algorithm to the Training set
#=====

from sklearn.preprocessing import PolynomialFeatures
polyagent = PolynomialFeatures(degree=5)
X_Poly = polyagent.fit_transform(X_train)
```

In [86]:

```
linReg1.fit (X_Poly, y_train )
```

Out[86]:

```
LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,
                  normalize=False)
```

In [87]:

```
X_Poly_test = polyagent.fit_transform(X_test)
```

In [88]:

```
score = linReg1.score(X_Poly,y_train)
print(score)
```

0.4927588433614648


In [89]:

```
score = linReg1.score(X_Poly_test,y_test)
print(score)
```

0.46653755542283765

In [90]:

```
#=====
# Visualize the poly regressor algo outcome
#=====
# Visualising the Regression results
#plt.scatter(X_Poly_test, y_test, color = 'red')
plt.plot(X_Poly_test, linReg1.predict(X_Poly_test), color = 'blue')
#plt.title('Median House Price Prediction')
plt.xlabel('Median Income')
plt.ylabel('Median House Price')
plt.show()
```



In [91]:

```
#Let's fit DT Reg

from sklearn.tree import DecisionTreeRegressor
DTRegressor = DecisionTreeRegressor(max_depth=3)
DTRegressor.fit(X_train,y_train)
```

Out[91]:

```
DecisionTreeRegressor(criterion='mse', max_depth=3, max_features=None,
                      max_leaf_nodes=None, min_impurity_decrease=0.0,
                      min_impurity_split=None, min_samples_leaf=1,
                      min_samples_split=2, min_weight_fraction_leaf=0.0,
                      presort=False, random_state=None, splitter='best')
```

In [92]:

```
#Quick check accuracy of the model on train
score = DTRegressor.score(X_train,y_train)
print(score)
```

0.48925102276025023

In [93]:

```
#=====
# Visualize the DT regressor algo outcome
#=====
# Visualising the Regression results
plt.scatter(X_test, y_test, color = 'red')
plt.plot(X_test, DTRegressor.predict(X_test), color = 'blue')
#plt.title('Median House Price Prediction')
plt.xlabel('Median Income')
plt.ylabel('Median House Price')
plt.show()
```

